

STAT430: Machine Learning for Financial Data

LarryHua.com/teaching

Spring 2019

Anatomy of a neural network

Core components of a neural network:

- Layers: which are combined into a network (or model)
- Input/features/covariates and targets/labels/response variables
- Loss function: the quantity that will be minimized during training
- Optimizer: a specific variant of stochastic gradient descent, determines how the model is updated based on the feedback from the loss function

Layers

- Some basic layers in Keras:
 - Simple vector data: 2D tensors of shape (samples, features) | densely connected layers
 - `layer_dense`: densely connected layers
 - Sequence data: 3D tensors of shape (samples, timesteps, features) | recurrent layers
 - `layer_lstm`: recurrent layers with long short-term memory
 - `layer_gru`: recurrent layers with gated recurrent unit
 - Image data: 4D tensors | 2D convolution layers
 - `layer_conv_2d`
- There are a lot more!
 - [Keras cheatsheet](#)

Layer compatibility

- Every layer only accepts input tensors of a certain shape and returns output tensors of a certain shape
- Example:

```
model <- keras_model_sequential() %>%  
  layer_dense(units = 32, input_shape = c(784)) %>%  
  layer_dense(units = 32)
```

- `units` and `input_shape` are the shapes w/o the sample axis for output and input, respectively
 - Only the first layer needs `input_shape`
- Pipe operator: `%>%`
 - From the `magrittr` package
 - Pass the value on its left as the **first argument** to the function on its right
 - `layer_dense(object, units, ...)`, object: model or layer object

Models, Loss functions

- A deep-learning model is a directed, acyclic graph of layers
- Picking the right network architecture is more an art than a science
- Multiple loss functions can be employed for multiple outputs, but losses have to be combined into a scalar value for gradient descents
 - Cross entropy are often used for a classification problem
 - `model %>% compile(loss = 'categorical_crossentropy', ...)`
 - `model %>% compile(loss = 'sparse_categorical_crossentropy', ...)`, when labels are integers and not one-hot encoding
 - MSE for a regression problem
 - `model %>% compile(loss = 'mse', ...)`

Develop models with Keras

- Two approaches:
 - Use `keras_model_sequential()`: a linear stacks of layers
 - Easier and more readable

```
model <- keras_model_sequential() %>%  
  layer_dense(units = 32, activation = "relu", input_shape = c(784)) %>%  
  layer_dense(units = 10, activation = "softmax")
```

- Use functional API: completely arbitrary architectures
- Much more flexible

```
input_tensor <- layer_input(shape = c(784))  
output_tensor <- input_tensor %>%  
  layer_dense(units = 32, activation = "relu") %>%  
  layer_dense(units = 10, activation = "softmax")  
model <- keras_model(inputs = input_tensor, outputs = output_tensor)
```

Compile and fit models with Keras

- Use `compile` to configure learning process

```
model %>% compile(  
  loss = 'categorical_crossentropy',  
  optimizer = optimizer_rmsprop(),  
  metrics = c('accuracy') # monitor performance  
)
```

- Use `fit` to configure run settings and fit model

```
his <- model %>% fit(  
  x_train, y_train,  
  epochs = 30, batch_size = 128, # epochs: number of passes through the full training set  
  validation_split = 0.2 # 20% from _train used for validation  
)
```

- [Try R](#)
- [Back to Course Scheduler](#)